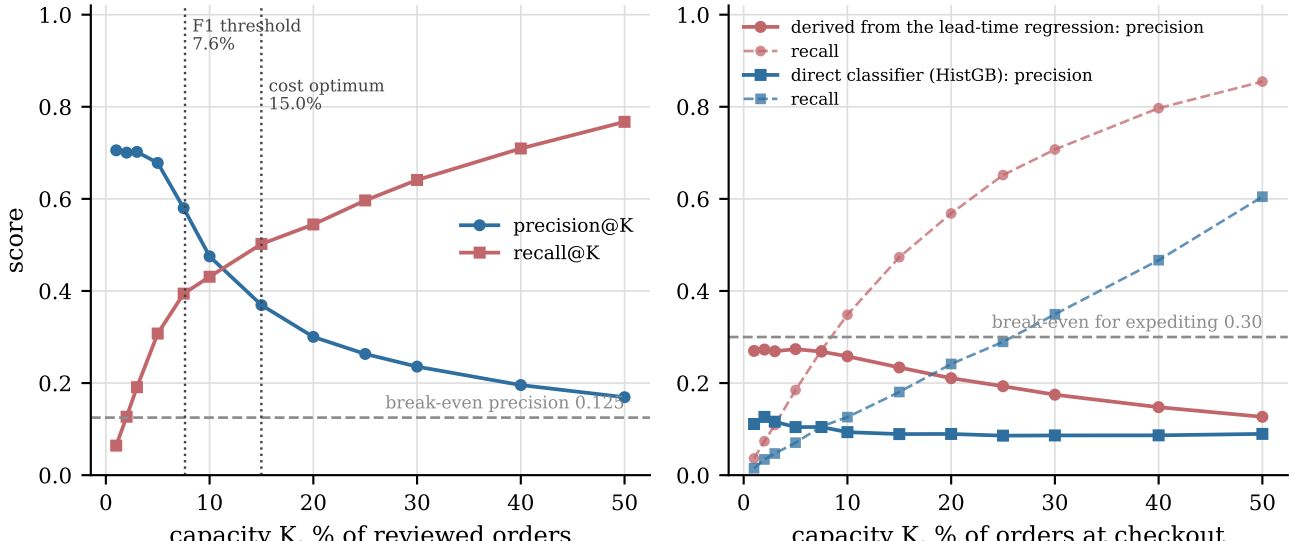
(a) Review risk: the top-K list
F1 threshold
7.6%
cost optimum
15.0%
precision@K
recall@K
break-even precision 0.125
score
capacity K, % of reviewed orders
(b) Late flag: derived against direct
derived from the lead-time regression: precision
recall
direct classifier (HistGB): precision
recall
break-even for expediting 0.30
capacity K, % of orders at checkout